

CMPS 101

Homework Assignment 6

Solutions

1. Show that any connected graph G satisfies $|E(G)| \geq |V(G)| - 1$. Hint: use induction on the number of edges.

Proof:

Let $m = |E(G)|$ and $n = |V(G)|$. We proceed by induction on m .

- I. Let $m = 0$. Then being connected, G can have only one vertex, hence $n = 1$. Therefore $m \geq n - 1$ reduces to $0 \geq 0$, and the base case is verified.
- II. Let $m > 0$. Assume for any connected graph G' with fewer than m edges: $|E(G')| \geq |V(G')| - 1$. We must show that $|E(G)| \geq |V(G)| - 1$, i.e. $m \geq n - 1$. Remove an edge $e \in E(G)$ and let $G - e$ denote the resulting subgraph. We have two cases to consider.

Case 1: $G - e$ is connected. Note that $G - e$ has n vertices and $m - 1$ edges, so the induction hypothesis gives $m - 1 \geq n - 1$. Certainly then $m \geq n - 1$, as was claimed.

Case 2: $G - e$ is disconnected. In this case $G - e$ consists of two connected components. (**See the claim and proof below.) Call them H_1 and H_2 . Observe that each component contains fewer than m edges. Suppose H_i has m_i edges and n_i vertices ($i = 1, 2$). The induction hypothesis gives $m_i \geq n_i - 1$ ($i = 1, 2$). Also $n = n_1 + n_2$ since no vertices were removed. Therefore

$$m = m_1 + m_2 + 1 \geq (n_1 - 1) + (n_2 - 1) + 1 = n_1 + n_2 - 1 = n - 1,$$


and therefore $m \geq n - 1$ as required. ///

(Note to grader: the following claim and proof are not necessary for full credit.)

Claim:** Let G be a connected graph and $e \in E(G)$, and suppose that $G - e$ is disconnected. (Such an edge e is called a *bridge*). Then $G - e$ has exactly two connected components.

Proof:

Since $G - e$ is disconnected, it has at least two components. We must show that it also has at most two components. Let e have end vertices u , and v . Let C_u and C_v be the connected components of $G - e$ that contain u and v respectively. Choose $x \in V(G)$ arbitrarily, and let P be an x - u path in G (note P exists since G is connected.) Either P includes the edge e , or it does not. If P does not contain e , then P remains intact after the removal of e , and hence P is an x - u path in $G - e$, whence $x \in C_u$. If on the other hand P does contain the edge e , then e must be the last edge along P from x to u .



In this case $P - e$ is an $x-v$ path in $G - e$, whence $x \in C_v$. Since x was arbitrary, every vertex in $G - e$ belongs to either C_u or C_v , and therefore $G - e$ has at most two connected components. ///

2. Let G be a graph on n vertices, A its adjacency matrix (as described in the Graph Theory handout), and let $d \geq 0$ be an integer. Show that the ij^{th} entry in A^d is the number of walks in G of length d from vertex i to vertex j . (Hint: use weak induction on d starting at $d = 0$.)

Proof:

- I. Observe that $A^0 = I$, the $n \times n$ identity matrix, which has a 1 in its i^{th} row and j^{th} column if $i = j$, and 0 in its i^{th} row and j^{th} column if $i \neq j$. Indeed there is exactly one walk of length 0 (the trivial walk consisting of no edges) from vertex i to vertex j if $i = j$ and zero such walks if $i \neq j$. The base case is therefore satisfied.
- II. Now let $d \geq 0$, and assume that $(A^d)_{ij}$ is the number of length d walks from i to j (for any i and j .) We must show that $(A^{d+1})_{ij}$ gives the number of length $d + 1$ walks from i to j . The definition of matrix multiplication says:

$$(A^{d+1})_{ij} = (A^d \cdot A)_{ij} = \sum_{k=1}^n (A^d)_{ik} \cdot A_{kj}$$

A walk from i to j of length $d + 1$ consists of two segments: (1) a walk of length d from i to some intermediate vertex k , followed by (2) the traversal of a single edge from k to j . Segment (1) can be constructed in $(A^d)_{ik}$ ways by the induction hypothesis. Segment (2) exists if and only if there is an edge from k to j , and can therefore be completed in A_{kj} ways (i.e. 1 way if there is an edge from k to j , and 0 ways otherwise). For a fixed intermediate vertex k the number of walks of length $d + 1$ from i to j is therefore $(A^d)_{ik} \cdot A_{kj}$. Summing over all possible k gives the total number of walks from i to j of length $d + 1$. By the above equation this number is $(A^{d+1})_{ij}$. The result now follows for all d by induction. ///

3. Use the result of the preceding problem to prove the following. The distance from vertex i to vertex j is equal to $\min\{d \mid (A^d)_{ij} > 0\}$. Here $(A^d)_{ij}$ denotes the ij^{th} entry in the matrix A^d .


Solution:

First note that a shortest i - j walk in a graph is actually a shortest i - j path, since if a walk visits a vertex x more than once, it can be made shorter by eliminating some edge traversals that constitute a closed walk starting and stopping at x . Thus our problem is to show that $\min\{d \mid (A^d)_{ij} > 0\}$ is the length of a shortest i - j walk.

By the result of the preceding problem, $(A^d)_{ij}$ is zero if and only if there is no i - j walk of length d in the graph. Consider the integer sequence: $(A^0)_{ij}$, $(A^1)_{ij}$, $(A^2)_{ij}$, $(A^3)_{ij}$, \dots . The first non-zero term in this sequence must therefore indicate the existence of at least one i - j walk, whose length is the corresponding exponent on A . Being the first such term, there can be no shorter i - j walks in the graph. Hence $\min\{d \mid (A^d)_{ij} > 0\}$ is the length of a shortest i - j walk in the graph. ■

4. p. 601: 22.2-2

Run BFS on the graph below using the following vertices as source. For each source, show the d and p values of each vertex, the order in which vertices are added to the Queue, and the state of the BFS tree after execution completes. Assume adjacency lists are processed in alphabetical order.




a. Let vertex u be the source

Solution:

vertex	adjacency list	d	p
r	s v	4	s
s	r w	3	w
t	u w x	1	u
u	t x y	0	nil
v	r	5	r
w	s t x	2	t
x	t u w y	1	u
y	u x	1	u

Queue: u t x y w s r v




b. Let vertex w be the source

Solution:

vertex	adjacency list	d	p
r	s v	2	s
s	r w	1	w
t	u w x	1	w
u	t x y	2	t
v	r	3	r
w	s t x	0	nil
x	t u w y	1	w
y	u x	2	x

Queue: w s t x r u y v




- c. Let vertex v be the source

Solution:

vertex	adjacency list	d	p
r	s v	1	v
s	r w	2	r
t	u w x	4	w
u	t x y	5	t
v	r	0	nil
w	s t x	3	s
x	t u w y	4	w
y	u x	5	x

Queue: v r s w t x u y



5. p. 602: 22.2-4

What is the run time of BFS if we represent its input graph by an adjacency matrix and modify the algorithm to handle this form of input?

Solution:

We re-write BFS so that the input graph is represented by an adjacency matrix, i.e. a 2-dimensional array A of ints, where $A[x][y] == 1$ if there is an edge from vertex x to vertex y , and $A[x][y] == 0$ otherwise.

```
BFS(G, s)
1.  for all x in V-{s}
2.      color[x] = white
3.      d[x] = infinity
4.      p[x] = nil
5.  color[s] = gray
6.  d[s] = 0
7.  p[s] = nil
8.  Q = new empty Queue
9.  Enqueue(Q, s)
10. while Q is not empty
11.     x = Dequeue(Q)
12.     for y = 1 to |V|
13.         if A[x][y]==1 and color[y]==white
14.             color[y] = gray
15.             d[y] = d[x] + 1
16.             p[y] = x
17.             Enqueue(Q, y)
18.     color[x] = black
```

Let $n = |V|$ and $m = |E|$. Observe that initialization (lines 1-9) has cost $\Theta(n)$, as it did in the earlier version of BFS. The total cost of Queue operations (lines 11 and 17) is also $\Theta(n)$ as before, since every vertex enters and exits the Queue exactly once (in worst case). The while loop (lines 10-18) executes n times (in worst case), and the inner for loop (lines 12-17) executes n times for every iteration of the outer while loop. Hence the for loop body executes a total of n^2 times, and the cost of the while loop is therefore $\Theta(n^2)$. Putting these together, the total cost of this version of BFS is $\Theta(n + n + n^2) = \Theta(n^2)$.

Although the problem does not ask for this, we should compare this run time to the $\Theta(n + m)$ run time for the adjacency list version of BFS. Note that in any graph or digraph with n vertices and m edges, we have $m = O(n^2)$. (In fact $m \leq n(n-1)/2$ in a graph, and $m \leq n^2$ in a digraph.) Thus the adjacency list run time $\Theta(n + m)$ can be no worse than the adjacency matrix $\Theta(n^2)$ run time. It can be much better though. If the graph (or digraph) is sparse, i.e. has many fewer than the maximum possible number of edges, then m is much less than n^2 , and $\Theta(n + m)$ is better than $\Theta(n^2)$. ///

6. p. 602: 22.2-7

There are two types of professional wrestlers: “good guys” and “bad guys.” Between any pair of professional wrestlers, there may or may not be a rivalry. Suppose we have n professional wrestlers and we have a list of r pairs of wrestlers for which there are rivalries. Give an $O(n+r)$ -time algorithm that determines whether it is possible to designate some of the wrestlers as good guys and the remainder as bad guys such that each rivalry is between a good guy and a bad guy. If it is possible to perform such a designation, your algorithm should produce it. (Hint: figure out how to use BFS to solve this problem.)

Solution:

Let the n wrestlers be the vertices of a graph in which two vertices are joined by an edge if and only if there is a rivalry between the corresponding wrestlers. The question asks whether it is possible to partition the vertex set into two subsets *Good* and *Bad* such that all edges join Good vertices to Bad vertices. Such a graph is called *bipartite*. More precisely, a graph $G = (V, E)$ is called bipartite if its vertex set V can be partitioned into two subsets X and Y in such a way that every edge has one end in X and the other end in Y . In other words, $\{u, v\} \in E$ implies that either $u \in X$ and $v \in Y$, or $u \in Y$ and $v \in X$. The pair of sets (X, Y) is called a *bipartition* of the vertex set V (See appendix B.4). Note that if $(v_0, v_1, v_2, \dots, v_k)$ is any path in a bipartite graph, then the vertices v_i ($i = 0, \dots, k$) must lie alternately in X and Y . Therefore any path starting in X and ending in Y must have odd length, while any path from X to X , or Y to Y must have even length.

Let s be any vertex in a connected graph G , and consider the distance fields $d[u]$ for all $u \in V$ after $\text{BFS}(G, s)$ has been run. If G is bipartite, then its vertex set V admits a bipartition (X, Y) as in the preceding paragraph, and we may assume for definiteness that $s \in X$. It follows from the above discussion that every vertex in X has an even distance from s , and every vertex in Y has an odd distance from s . (What is true for all paths must also be true for shortest paths.) Thus if G is bipartite, the bipartition must be given by

$$X = \{ u \in V : d[u] \text{ is even} \} \quad \text{and} \quad Y = \{ u \in V : d[u] \text{ is odd} \}.$$

Conversely, if G is not bipartite, then no bipartition of $V(G)$ exists. Thus if we define X and Y as above, there must either exist an edge joining X to X , or an edge joining Y to Y . We have proved the following: G is bipartite if and only if no two adjacent vertices have the same mod 2 distance from s . The following algorithm exploits this fact. However it works only in the special case that G is a connected graph.

isBipartite(G) (Pre: G is a connected graph)

1. pick any $s \in V[G]$
2. $\text{BFS}(G, s)$
3. for all $u \in V[G]$
4. for all $v \in \text{adj}[u]$
5. if $d[u] \equiv d[v] \pmod{2}$
6. print “no bipartition of wrestlers is possible”
7. return
8. print “The good guys are:”
9. print the set of vertices that are an even distance from s

The cost of this algorithm is obviously $O(|V| + |E|)$ in worst case. Of course, there is no reason to expect that the graph corresponding to wrestler rivalries would be connected. A moments thought

shows that a disconnected graph is bipartite if and only if all of its connected components are bipartite. In the modified algorithm below, T and X denote subsets of $V(G)$, and $G - T$ denotes the subgraph of G obtained by removing all vertices in T , along with all incident edges.


IsBipartite(G)

1. $T = X = \emptyset$
2. while $V - T \neq \emptyset$
3. pick $s \in V - T$
4. BFS($G - T, s$)
5. $T = T \cup \{\text{black vertices}\}$
6. $X = X \cup \{\text{black vertices whose distance from } s \text{ is even}\}$
7. for all black vertices u
8. for all $v \in \text{adj}[u]$
9. if $d[u] \equiv d[v] \pmod{2}$
10. print “no such partition of the wrestlers is possible”
11. return
12. print “the good guys are:”
13. print the members of X

7. p.610: 22.3-2


Run DFS on the graph of Figure 22.6 (p.611), shown below. Assume the main loop of DFS (lines 5-7) processes vertices in alphabetical order, and assume that each adjacency list is also processed in alphabetical order. Show the discover and finishing times for each vertex, and show the classification of each edge.

Solution:



Vertex	Adj	Discover	Finish
q	s t w	1	16
r	u y	17	20
s	v	2	7
t	x y	8	15
u	y	18	19
v	w	3	6
w	s	4	5
x	z	9	12
y	q	13	14
z	x	10	11

DFS Forest



Edge	Classification
(q, s)	tree
(q, t)	tree
(q, w)	forward
(r, u)	tree
(r, y)	cross
(s, v)	tree
(t, x)	tree
(t, y)	tree
(u, y)	cross
(v, w)	tree
(w, s)	back
(x, z)	tree
(y, q)	back
(z, x)	back